

ROIL: Robust Offline Imitation Learning

Gersi Doko¹, Guang Yang², Daniel S. Brown² Marek Petrik¹

¹University of New Hampshire, ²University of Utah

Summary

Motivation

- ▶ Compute policies that are robust to parameter uncertainty is very important in many domains, like health care, inventory control or finance.
- ▶ Seek policies that maximize the expected return over a distribution of MDP models

Limitations of existing methods

- ▶ Reliance on \hat{u}_e leads to covariate shift for off-policy datasets.
- ▶ Inability to specify reliance on \hat{u}_e .
- ▶ No guarantees of policy convergence to u_e .


Our contributions

- ▶ New algorithm for robust offline imitation learning.
- ▶ Guaranteed convergence to the optimal policy for tabular domains.
- ▶ Flexibility to define the reliance on \hat{u}_e .

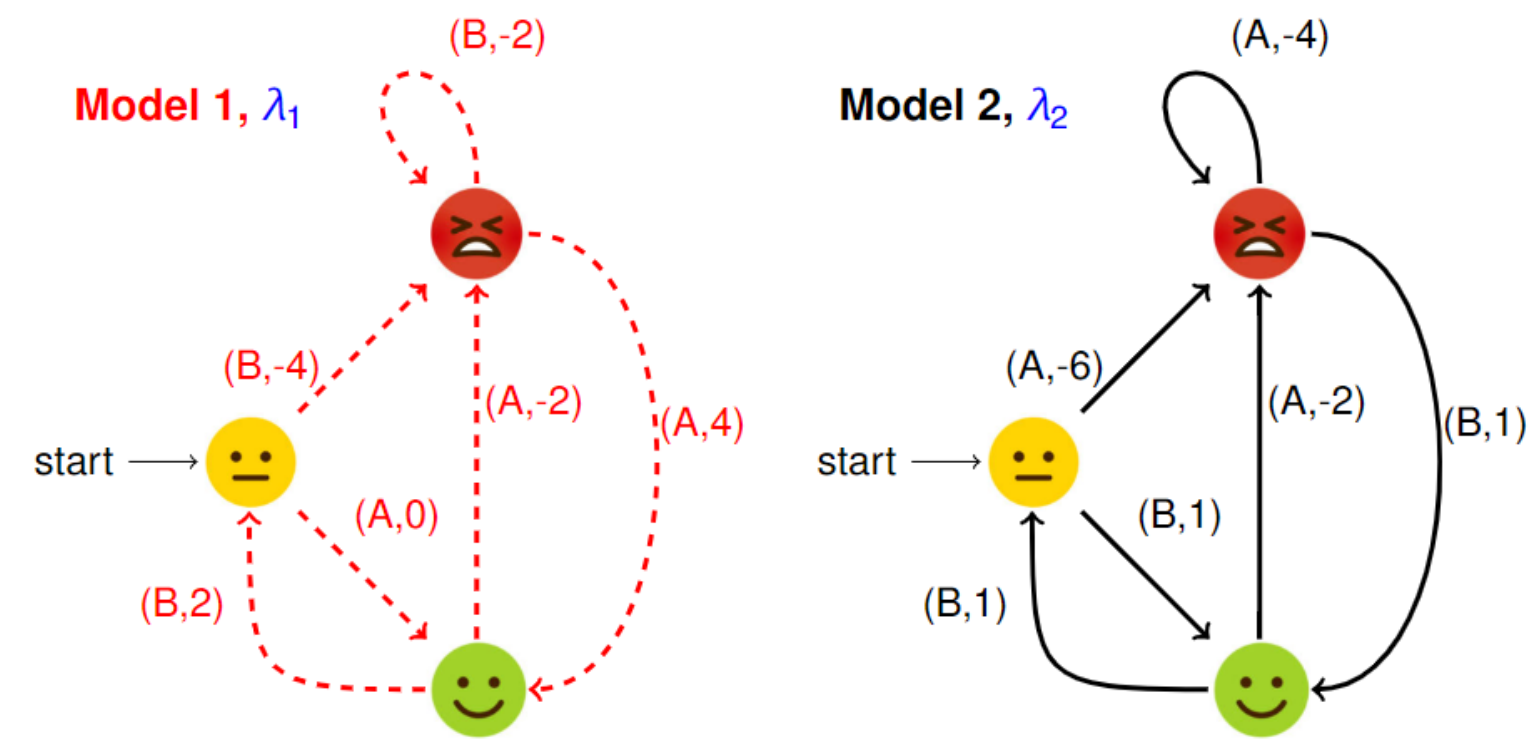
Markov Decision Process (MDP)

MDP : consists of a tuple $\langle \mathcal{S}, \mathcal{A}, p, r \rangle$

- State space $\mathcal{S} = \{ \text{mild, moderate, severe} \}$
- Action space $\mathcal{A} = \{A, B\}$
- Transition probability p : $p(\text{mild}, \text{mild}, A) = 1$
- Reward function r : $r(\text{mild}, A) = 100$



Multi-model Markov Decision Processes (MMDPs)



$\mathcal{M} = \{1, 2\}, \lambda = \{\lambda_1, \lambda_2\}, \mathcal{S} = \{ \text{mild, moderate, severe} \}, \mathcal{A} = \{A, B\}$

- ▶ Mean return across the uncertain true models

$$\rho(\pi) = \mathbb{E}^{\lambda} \left[\mathbb{E}^{\pi, p^{\tilde{m}}, \mu} \left[\sum_{t=1}^T r_t^{\tilde{m}}(\tilde{s}_t, \tilde{a}_t) \mid \tilde{m} \right] \right] \quad (1).$$

- ▶ Optimal policy ρ^*

$$\rho^* = \max_{\pi \in \Pi} \rho(\pi).$$

Prior Work: Weight-Select-Update (WSU)

WSU Approximation Algorithm

Input: MMDPs, Model weights λ

Output: $\pi = (\pi_1, \dots, \pi_T)$

1. Initialize $v_{T+1, m}^{\pi}(s_{T+1}) = 0, \forall m \in \mathcal{M}$
2. For $t = T, T-1, \dots, 1$ do
3. $\pi_t(s_t) \in \arg \max_{a \in \mathcal{A}} \sum_{m \in \mathcal{M}} \lambda_m \cdot q_{t, m}^{\pi}(s_t, a), \quad \forall s_t \in \mathcal{S}.$
4. $v_{t, m}^{\pi}(s_t) = r_t^m(s_t, \pi(s_t)) + \sum_{s_{t+1} \in \mathcal{S}} p_t^m(s_{t+1} \mid s_t, \pi(s_t)) \cdot v_{t+1, m}^{\pi}(s_{t+1}), \forall m \in \mathcal{M}.$
5. end for

MMDP Policy Gradient

- ▶ **Main idea:** Take a coordinate ascent perspective to adjust model weights iteratively.
- ▶ **Definition 4.1** An *adjustable weight* for each $m \in \mathcal{M}, \pi \in \Pi, t \in \mathcal{T}$, and $s \in \mathcal{S}$ is

$$b_{t, m}^{\pi}(s) = \mathbb{P}[\tilde{m} = m, \tilde{s}_t = s],$$

where $S_0 \sim \mu, \tilde{m} \sim \lambda$, and $\tilde{s}_1, \dots, \tilde{s}_T$ are distributed according to $p^{\tilde{m}}$ of policy π .

- ▶ **Theorem 4.1:** Gradient of ρ in Eq. (1) for each $t \in \mathcal{T}, \hat{s} \in \mathcal{S}, \hat{a} \in \mathcal{A}$, and $\pi \in \Pi_R$ is

$$\frac{\partial \rho(\pi)}{\partial \pi_t(\hat{s}, \hat{a})} = \sum_{m \in \mathcal{M}} b_{t, m}^{\pi}(\hat{s}) \cdot q_{t, m}^{\pi}(\hat{s}, \hat{a}),$$

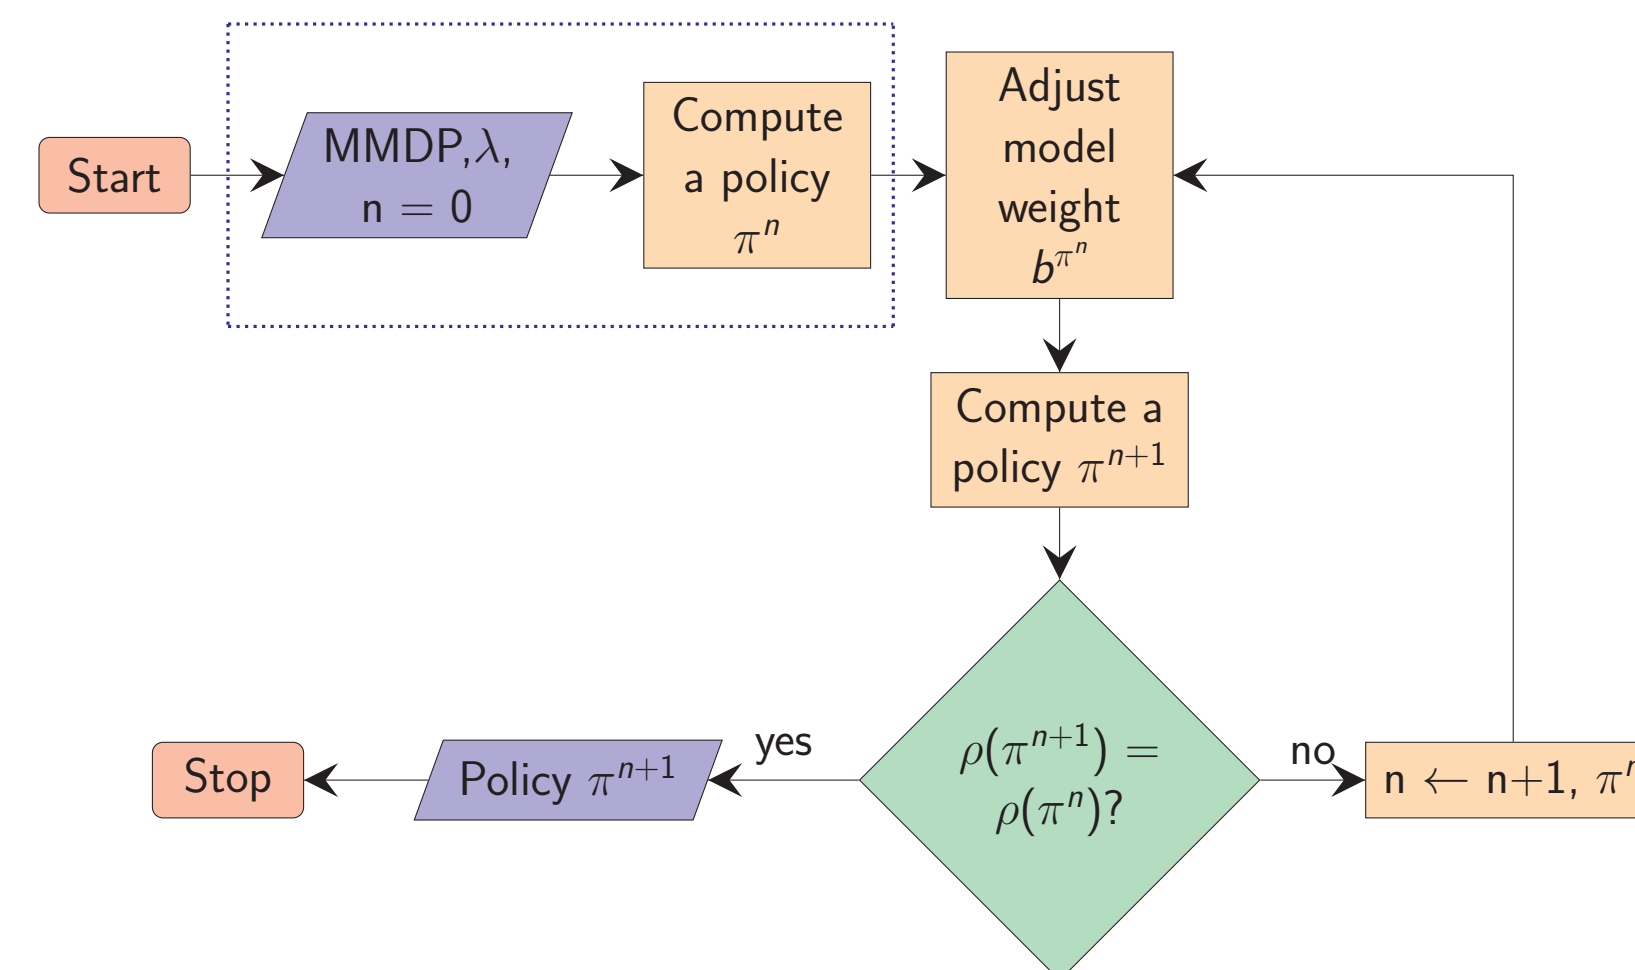
where q is state-action value function and b is an adjustable weight

- ▶ **Corollary 4.2** For any $\bar{\pi} \in \Pi$ and $t \in \mathcal{T}$, function $\pi_t \mapsto \rho(\bar{\pi}_1, \dots, \pi_t, \dots, \bar{\pi}_T)$ is *linear*.
- ▶ Linearity implies that we can solve the maximization over $\pi_t(s)$ as

$$\pi_t^n(s) \in \arg \max_{a \in \mathcal{A}} \sum_{m \in \mathcal{M}} b_{t, m}^{\pi^{n-1}}(s) \cdot q_{t, m}^{\pi^n}(s, a).$$

Coordinate Ascent Dynamic Programming (CADP)

- ▶ **Main idea:** Combine coordinate ascent method and DP to solve MMDPs.
- ▶ **Corresponds to:** Replace the fixed model weights λ_m in WSU by adjustable weights $b_{t, m}^{\pi}$
- ▶ Blue dotted rectangle is to compute an initial policy (for example by WSU, MVP)




Related Algorithms


- ▶ Prior MMDP algorithms: WSU and MVP
- ▶ Gradient-based MMDP methods: Mirror and Gradient
- ▶ Thompson sampling-based algorithms: MixTS
- ▶ POMDP formulations: QMDP and POMCP

Simulation Results: Pest Control

- ▶ Time horizon $T = 50$, Domain: Pest control simulation
- ▶ Below figure: mean returns of CADP with different initial policies.



- ▶ Left figure: mean returns of algorithms, and right figure: runtimes of algorithms.
- ▶ Marker X: no single policy available or runtime is greater than 900 minutes



Simulation Results: Other Domains

- ▶ Mean returns $\rho(\pi)$ on the test set of policies π computed by each algorithm

Algorithm	RS		POP		POPS		INV		HIV	
	T = 50	T = 150	T = 50	T = 150	T = 50	T = 150	T = 50	T = 150	T = 5	T = 20
CADP	204	207	-361	-368	-1067	-1082	323	350	33348	42566
WSU	203	206	-542	-551	-1915	-1932	323	349	33348	42564
MVP	201	204	-704	-717	-2147	-2179	323	350	33348	42564
Mirror	181	183	-1650	-1600	-3676	-3800	314	345	33348	42566
Gradient	203	206	-542	-551	-1915	-1932	323	349	33348	42564
MixTS	167	176	-1761	-1711	-2857	-3016	327	350	293	-1026
QMDP	190	183	-	-	-	-	-	-	30705	39626
POMCP	58	64	-	-	-	-	-	-	25794	30910
Oracle	210	213	-168	-172	-882	-894	332	360	40159	53856