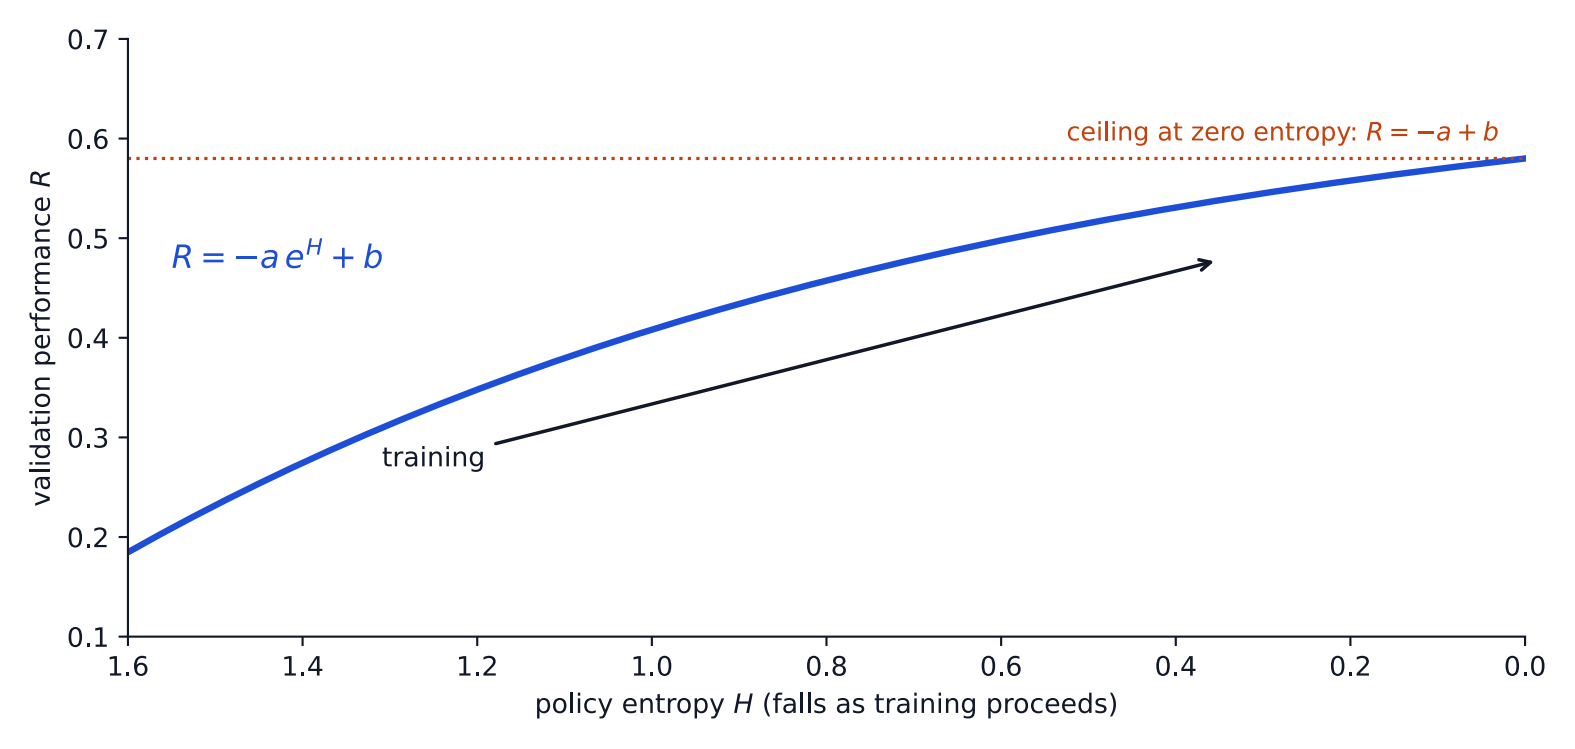
0.7
0.6
0.5
0.4
0.3
0.2
0.1
validation performance R
ceiling at zero entropy: $R = -a + b$
$R = -a\,e^{H} + b$
training
1.6
1.4
1.2
1.0
0.8
0.6
0.4
0.2
0.0
policy entropy H (falls as training proceeds)